

Adding API Keys, Smart Routing & Free Fallback

Prompts and quick reference for this lesson

This PDF contains all the prompts from the lesson so you can copy and paste them directly instead of typing them out. It also includes both versions of the routing rules (Anthropic + OpenAI and OpenAI-only) plus the environment variable names and useful links.

Running OpenClaw locally? In the video I add API keys through the Hostinger Docker Manager. If you're on a local install, just paste your key directly to your agent on Telegram and ask it to add it to your config. See the "Note for Local Install Users" PDF at the start of the course for details.

Useful Links

- OpenClaw OAuth docs: docs.openclaw.ai/concepts/oauth
- OpenAI API platform: platform.openai.com
- NVIDIA build (Kimi K2.5 free key): build.nvidia.com
- Anthropic console: console.anthropic.com

Environment Variable Names

Provider	Variable Name	How You Pay
OpenAI (API key)	OPENAI_API_KEY	Pay per token
OpenAI (ChatGPT subscription)	Connected via OAuth	Flat monthly fee
Anthropic (Claude)	ANTHROPIC_API_KEY	Pay per token
NVIDIA (Kimi K2.5 free)	MOONSHOT_API_KEY	Free

Prompts Used in the Lesson

Connect ChatGPT subscription via OAuth

Read the OpenClaw OAuth documentation at <https://docs.openclaw.ai/concepts/oauth> and then help me connect my ChatGPT subscription to OpenClaw through Codex OAuth so I can use my existing plan instead of paying per token.

Yes, go ahead and run the Codex OAuth login command for me.

Keep Claude as the default. We're going to set up routing rules in a moment that will tell you when to use which model.

Clean up unused fallback models

Show me your current model configuration and remove any fallback models that we don't have API keys or subscriptions for. Only keep the ones that are actually connected.

Add Kimi K2.5 as free fallback

I've added a free Kimi K2.5 API key through NVIDIA as an environment variable called MOONSHOT_API_KEY. Set it up as my very last fallback model, only to be used when all other models are unavailable.

Routing Rules

Pick the version that matches your setup. If you have both Anthropic and OpenAI connected, use Version A. If you only have a ChatGPT subscription without Anthropic, use Version B.

Version A — Anthropic + OpenAI (recommended)

For students who have both an Anthropic API key and a ChatGPT subscription connected:

Here are my model routing rules. Save these as a permanent rule and update your configuration.

Use Claude Sonnet as the default model for everyday tasks. If Sonnet isn't available, fall back to Codex from my ChatGPT subscription.

For any tasks that need serious thinking and planning, use Claude Opus. If Opus isn't available, fall back to the best available model.

For coding tasks, use Codex from my ChatGPT subscription.

For simple routine tasks like quick lookups and heartbeat runs, use Claude Haiku. If Haiku isn't available, fall back to Codex.

When tackling a big task, use Opus to come up with the plan, but then hand off the actual work to cheaper models whenever possible.

Every time you run a task, tell me which model you're using so I can check that the routing is working.

Version B — OpenAI subscription only

For students who only have a ChatGPT subscription without Anthropic:

Here are my model routing rules. Save these as a permanent rule and update your configuration.

Use Codex as my default model for everyday tasks and coding.

For simple routine tasks like quick lookups and heartbeat runs, switch to the cheapest available model to save my subscription quota.

When tackling a big task, spend the time thinking it through properly, but then hand off the repetitive parts to cheaper models whenever possible.

Every time you run a task, tell me which model you're using so I can check that the routing is working.

Understanding Model Tags in the Dashboard

When you look at the model dropdown in your gateway dashboard, each model has a tag on the right side. This tag tells you which provider and billing method is being used — not the model name on the left.

Tag	What It Means	How You Pay
openai-codex	ChatGPT subscription (OAuth)	Covered by your monthly fee
openai	OpenAI API key	Per token from your credit balance
anthropic	Anthropic API key	Per token from your credit balance
moonshot	NVIDIA free API (Kimi)	Free (rate limited)

Quick rule: If you want to use your ChatGPT subscription without paying anything extra, always pick the model tagged *openai-codex*. A model name containing the word "codex" doesn't necessarily mean it goes through your subscription — only the tag on the right tells you that.

Kimi K2.5 Free Fallback — Quick Setup Steps

1. Go to build.nvidia.com and search for Kimi K2.5
2. Create a free NVIDIA account if you don't have one
3. Click View Code on the Kimi K2.5 page
4. Click Generate API Key and copy the full key (starts with NVAPI)
5. Add it as an environment variable called MOONSHOT_API_KEY
6. VPS: Hostinger Docker Manager → Environment → Save and Deploy
7. Local: paste the key to your agent and ask it to add MOONSHOT_API_KEY to the config
8. Tell your agent to set it up as the last fallback (prompt above)

About the free version: Kimi through NVIDIA is slower than paid models and has rate limits. That's normal for a free service. Its job here is to act as a safety net so your agent can still let you know what's going on if your main models run out of credits or hit an outage.

A Note About Anthropic Subscriptions

Unlike OpenAI, Anthropic does not officially allow connecting a Claude subscription (Pro/Team) to third-party tools like OpenClaw. Some users have had their accounts suspended for doing this. The only supported way to use Claude models with OpenClaw is through an Anthropic API key with pay-per-token credits, which is what we set up earlier in the course.