

OpenClaw AI Models Guide

Your complete reference for free, cheap, and premium AI models in OpenClaw

This guide brings together everything covered in the Models & Cost Control section of the course. Use it as a quick reference when you need to look up pricing, environment variable names, or setup steps.

Section	
1. How AI Model Pricing Works	Tokens, API keys vs subscriptions
2. Model Tiers	Tier 1 (premium) to free
3. All Your Options at a Glance	Quick overview of every option
4. Running OpenClaw for Free or Cheap	OpenRouter, Gemini, Moonshot, Mini
5. Connecting Paid Models	Anthropic API, ChatGPT OAuth, OpenAI API
6. Smart Routing Rules	Ready-to-use prompts
7. Free Fallback: Kimi & Ollama	NVIDIA setup, local models
8. Recommended Setups by Budget	\$0 to \$200/month
9. Quick Reference: All Models	Prices, context, best use
10. Environment Variable Names	Copy-paste ready
11. Useful Links	All provider URLs

1. How AI Model Pricing Works

OpenClaw itself is completely free and open source. What costs money is the AI model that powers it. There are two different ways to pay:

API Key (Pay Per Use)	Subscription (Flat Rate via OAuth)
Create an API key, load credits, pay per token used. Like a pay-as-you-go phone plan.	Connect your existing ChatGPT subscription to OpenClaw via OAuth. No extra per-token cost — covered by your monthly plan.
Works with: Anthropic, OpenAI, Google, Moonshot, OpenRouter, and many others	Works with: OpenAI ChatGPT Plus (\$20/mo) or Pro (\$200/mo)

What are tokens? Tokens are the units AI models use to process text. Roughly 1 token = 3/4 of a word, so 1,000 tokens is about 750 words. When you message your agent, it costs input tokens. When the agent replies, it costs output tokens. OpenClaw also sends your workspace files, memory, and conversation history with every message — that's why a simple question can use 50,000–100,000 tokens before the model even starts thinking.

2. Model Tiers

AI models come in different tiers. The general rule is: use a strong model for thinking and planning, and a cheap model for routine tasks.

Tier	Examples	Cost Range	Best For
Tier 1 Premium	Claude Opus GPT-5.4 Pro	\$15–75/M output	Complex reasoning, planning
Tier 2 Daily Driver	Claude Sonnet GPT-5.4	\$3–15/M output	Everyday tasks, coding
Tier 3 Budget	Claude Haiku GPT-5.4 Mini	\$0.25–4/M output	Routine tasks, heartbeats
Free	MiniMax M2.5 (OpenRouter) Gemini Flash	\$0	Budget setups, fallbacks

Haiku is 25x cheaper than Opus. A single Opus prompt can cost \$2–6 if it's doing a lot of work. That's why smart routing matters.

3. All Your Options at a Glance

Here's a quick overview of every option available to you. You can mix and match any of these — OpenClaw doesn't lock you into one provider. Each option is covered in detail in the sections that follow.

Free Models

- **MiniMax M2.5 via OpenRouter** — free, solid quality, good for everyday tasks
- **NVIDIA Nemotron 3 Super via OpenRouter** — free, best open model benchmark score for OpenClaw
- **Google Gemini Flash** — massive free tier (1,500 req/day, 1M tokens/min)
- **Kimi K2.5 via NVIDIA NIM** — free but slow (30+ sec), best as emergency fallback
- **Local models via Ollama** — \$0 per token, needs 16GB+ RAM VPS

Cheap Models (Under \$5/month)

- **Kimi K2.5 via OpenRouter** — \$0.42/M input, very popular in the OpenClaw community, reliable
- **Paid Kimi K2.5 via Moonshot** — \$0.60/M input, 5x cheaper than Sonnet, fast and reliable
- **OpenAI GPT-5.4 Mini** — \$0.25/M input, great for routine tasks (needs OpenAI API key)
- **Claude Haiku** — \$0.80/M input, cheapest Anthropic option

Premium Models

- **Claude Sonnet 4.6** — \$3/M input, recommended daily driver, best tool calling
- **Claude Opus 4.6** — \$15/M input, highest quality, use only for hard tasks
- **ChatGPT subscription via OAuth** — flat rate (\$20 or \$200/mo), GPT-5.4 included, OpenAI supports this
- **Anthropic subscription** — some people use it with OpenClaw without issues, but officially not allowed by Anthropic. Risk of account ban. Your choice — safest option is API key instead

Bottom line: OpenClaw works best with a paid model, but you can absolutely run it for free. The best approach for most people is to mix and match — use a cheap or free model for everyday tasks, and a premium model only when you really need it.

4. Running OpenClaw for Free or Cheap

You can technically run OpenClaw entirely on free models. OpenRouter has several free models like MiniMax M2.5 and NVIDIA Nemotron 3 Super that you can set as your primary. But free models have tight rate limits, they tend to be slower, and they can change without warning. It's possible, but it's not going to be the smoothest experience. A lot of people in the community use Kimi K2.5 through OpenRouter which isn't free but is very cheap at about \$0.42/M input tokens.

Free Models via OpenRouter — Best Free Primary Options

OpenRouter has several completely free models. MiniMax M2.5 and NVIDIA Nemotron 3 Super are both solid choices. Free models can change over time — check OpenRouter for the latest free options.

- Setup: Create free account at openrouter.ai → generate API key → add as `OPENROUTER_API_KEY`
- Tell agent: set `minimax/minimax-m2.5` (or `nvidia/nemotron-3-super`) as primary
- Switch between free models anytime — same API key, just change model identifier

Kimi K2.5 via OpenRouter — Very Cheap, Community Favourite

Not free, but extremely popular in the OpenClaw community. Many people use it as their main everyday model.

- Input: \$0.42/M | Output: \$2.20/M | Context: 262K
- Uses same `OPENROUTER_API_KEY` — just switch model identifier to `moonshotai/kimi-k2.5`
- Typical cost: \$2–5/month for daily use

Google Gemini Flash — Most Generous Free Tier

1,500 requests/day, 1 million tokens/minute. More than enough for daily use.

- Cost: Free (generous limits) | Context: 1M tokens
- Setup: Get API key at aistudio.google.com → add as `GEMINI_API_KEY`
- Note: Data may be used for training on free plan

Paid Kimi K2.5 via Moonshot — Under \$5/month

Much faster than the free NVIDIA version. Roughly 5x cheaper than Claude Sonnet. Many people use this as their main model.

- Input: \$0.60/M | Output: \$3.00/M | Context: 262K
- Setup: Create account at platform.moonshot.ai → top up → generate key → add as `MOONSHOTAI_API_KEY`
- Not a built-in provider — tell agent to configure custom provider with base URL `https://api.moonshot.cn/v1`

OpenAI GPT-5.4 Mini — Cheap Routine Tasks

If you already have an OpenAI API key, the Mini models are extremely cheap for heartbeat runs, lookups, and basic questions.

- Input: \$0.25/M | Output: \$2.00/M | Context: 272K
- Requires OPENAI_API_KEY (separate from ChatGPT subscription)
- Tell agent: *"For routine tasks and heartbeat runs, use GPT-5.4 Mini instead of more expensive models."*

5. Connecting Paid Models

Anthropic API Key (Claude Models)

This is what we use throughout the course. Go to console.anthropic.com, create an account, add at least \$40 in credits (to reach Tier 2 for smoother rate limits), set a monthly spend limit, and generate an API key. Add it as `ANTHROPIC_API_KEY` in your Hostinger Docker Manager.

ChatGPT Subscription via OAuth (OpenAI Codex)

If you have ChatGPT Plus (\$20/mo) or Pro (\$200/mo), connect it to OpenClaw through OAuth at no extra cost. Your agent walks you through the process — you get a URL, authorize in your browser, and paste the redirect URL back. You get GPT-5.4 through the `openai-codex` tag. OpenAI actively supports this.

Anthropic Subscription (Claude Pro/Max) — Use with Caution

Some people in the community connect their Claude Pro or Max subscription to OpenClaw, and it works without any issues. However, Anthropic has officially said that using your Claude subscription in third-party tools like OpenClaw is not allowed under their terms of service. Earlier this year they started enforcing this and some users had their accounts shut down. That said, plenty of people are still doing it without any problems. It's your choice if you want to take that risk — just be aware that if your account gets banned, you'd need to create a new one and any credits would be gone. The safest option with Anthropic is to use an API key with pay-per-use credits.

OpenAI API Key (Pay Per Use)

If you don't have a ChatGPT subscription, you can use the OpenAI API on a pay-per-use basis. Go to platform.openai.com (separate from chatgpt.com), create an account, add billing, generate an API key. Add as `OPENAI_API_KEY`. This gives you access to all GPT models including the cheap Mini variants.

Important: The OpenAI API platform (platform.openai.com) and your ChatGPT subscription (chatgpt.com) are completely separate. Having a ChatGPT account does NOT give you API access. Models tagged **openai** = API key (pay per use). Models tagged **openai-codex** = ChatGPT subscription (flat rate).

6. Smart Routing Rules

Smart routing tells your agent which model to use for which task. This can cut costs by 60–70%. Copy and paste these prompts into Telegram. Add "Save these as a permanent rule and update your configuration" at the end.

Option A: Anthropic + ChatGPT Subscription

Use Claude Sonnet as the default model for everyday tasks. If Sonnet isn't available, fall back to Codex from my ChatGPT subscription. For tasks that need serious thinking and planning, use Claude Opus. If Opus isn't available, fall back to the best available model. For coding tasks, use Codex from my ChatGPT subscription. For simple routine tasks like quick lookups and heartbeat runs, use Claude Haiku. If Haiku isn't available, fall back to Codex. When tackling a big task, use Opus to come up with the plan, but hand off the actual work to cheaper models. Every time you run a task, tell me which model you're using.

Option B: ChatGPT Subscription Only

Use Codex as my default model for everyday tasks and coding. For simple routine tasks like quick lookups and heartbeat runs, switch to the cheapest available model to save my subscription quota. When tackling a big task, spend the time thinking it through properly, but hand off repetitive parts to cheaper models. Every time you run a task, tell me which model you're using.

Option C: Free Models Only

Use MiniMax M2.5 through OpenRouter as my primary model. Set Gemini Flash as my first fallback. Set Kimi K2.5 through NVIDIA as my last fallback. Remove any paid models from the fallback chain. Every time you run a task, tell me which model you're using.

7. Free Fallback: Kimi & Ollama

If your paid provider runs out of credits or has an outage, your agent just goes silent — no error, no explanation. A free fallback can at least tell you what went wrong.

Kimi K2.5 via NVIDIA NIM — Setup Steps

- Go to build.nvidia.com, create free account, search for Kimi K2.5
- Click View Code → Generate API Key (starts with NVAPI — copy the whole thing)
- Hostinger Docker Manager → Manage → Environment → add MOONSHOT_API_KEY
- Save and Deploy, wait for restart
- In Telegram: *"I've added a free Kimi K2.5 API key through NVIDIA as MOONSHOT_API_KEY. Set it up as my very last fallback model."*
- If agent doesn't recognise the model: *"Check docs.openclaw.ai/providers for the Kimi K2.5 model identifier."*
- Slow (30+ sec sometimes), 40 RPM limit — use as emergency only

Local Models via Ollama

Runs on your server, \$0 per token. Needs KVM 4+ (16GB RAM). Ask your agent to check your specs and recommend a model. Make sure it supports tool usage.

- Website: ollama.com
- Recommended: Qwen 3.5 27B (needs ~20GB) | Lighter: Qwen 3.5 9B (8GB)

Prompt: *"I'm thinking about running a local model through Ollama. Can you check my server specs and recommend a model that would work well? It needs to support tool usage."*

8. Recommended Setups by Budget

Budget	Primary	Fallback	Est. Cost
\$0/month	MiniMax M2.5 (OpenRouter)	Gemini Flash → Kimi NVIDIA	\$0
\$2–5/month	Kimi K2.5 (OpenRouter)	Free model + Gemini	\$2–5
\$20/month	ChatGPT sub (Codex)	Free model (OpenRouter)	\$20 flat
\$20–50/month	Claude Sonnet (API)	Codex + Haiku	\$20–50
\$50–200/month	Sonnet + Opus	Codex + Haiku + Kimi	\$50–200

You can mix and match all of these. Use Claude for the hard stuff, Kimi for everyday tasks, and Gemini as a free safety net. Switch anytime using /model or by updating your routing rules.

9. Quick Reference: All Models at a Glance

Model	Provider	In \$/M	Out \$/M	Context	Best For
MiniMax M2.5	OpenRouter	Free	Free	197K	Free primary
Nemotron Super 3	OpenRouter	Free	Free	262K	Free primary
Gemini Flash	Google	Free*	Free*	1M	Free backup
Kimi K2.5	NVIDIA NIM	Free	Free	200K	Emergency fallback
Kimi K2.5	OpenRouter	\$0.42	\$2.20	262K	Cheap primary
Kimi K2.5	Moonshot	\$0.60	\$3.00	262K	Cheap primary
GPT-5.4 Mini	OpenAI API	\$0.25	\$2.00	272K	Cheap routine
Claude Haiku	Anthropic	\$0.80	\$4.00	200K	Cheap routine
GPT-5.4 Codex	ChatGPT Sub	Incl.	Incl.	272K	Flat rate
Claude Sonnet	Anthropic	\$3.00	\$15.00	200K	Best daily driver
Claude Opus	Anthropic	\$15.00	\$75.00	200K	Complex reasoning

* Gemini Flash free tier: 1,500 req/day, 1M tokens/min. Data may be used for training.

Prices as of March 2026. Always check the provider's website for current pricing.

10. Environment Variable Names

Add these in Hostinger Docker Manager → Manage → Environment. Use exactly these names:

Provider	Variable Name	Where to Get Key
Anthropic	ANTHROPIC_API_KEY	console.anthropic.com
OpenAI (API)	OPENAI_API_KEY	platform.openai.com
OpenRouter	OPENROUTER_API_KEY	openrouter.ai
Google Gemini	GEMINI_API_KEY	aistudio.google.com
Kimi (NVIDIA free)	MOONSHOT_API_KEY	build.nvidia.com
Kimi (Moonshot paid)	MOONSHOTAI_API_KEY	platform.moonshot.ai
Brave Search	BRAVE_API_KEY	brave.com/search/api

ChatGPT subscription (OAuth) does not use an environment variable — it's connected through an OAuth flow with your agent.

11. Useful Links

- **OpenRouter (free & paid models)** — openrouter.ai
- **OpenRouter — OpenClaw Integration** — openrouter.ai/docs/guides/guides/openclaw-integration
- **Google AI Studio (Gemini API key)** — aistudio.google.com
- **NVIDIA NIM (Kimi K2.5 free)** — build.nvidia.com
- **Moonshot Console (paid Kimi)** — platform.moonshot.ai
- **Moonshot — OpenClaw Guide** — platform.moonshot.ai/docs/guide/use-kimi-in-openclaw
- **OpenAI API Platform** — platform.openai.com
- **Anthropic Console** — console.anthropic.com
- **Ollama (local models)** — ollama.com
- **OpenClaw Documentation** — docs.openclaw.ai
- **OpenClaw OAuth Docs** — docs.openclaw.ai/concepts/oauth
- **OpenClaw Model Providers** — docs.openclaw.ai/providers