

OpenClaw + Google Workspace

Security Guide — Read This Before Connecting

In the next lesson we are going to connect our OpenClaw agent to Google Workspace using the GOG skill. This gives your agent access to Gmail, Calendar, Drive, Contacts, Sheets, and Docs. It is incredibly powerful and it unlocks a huge range of automations. But before we do that, there are some important security considerations you need to understand.

Why You Need to Think About This First

Here is a simple way to think about this. Imagine you are bringing someone new into your business. They are brilliant at what they do, they work around the clock, and they never complain. But it is their first week on the job. You do not know them that well yet. You would not hand them the keys to your personal email, your bank login, your private files, and your entire client list on day one. You would start them off with their own workspace, their own accounts, and limited access until you build up that trust over time.

That is exactly how you should think about connecting OpenClaw to Google Workspace. Your agent is incredibly capable, but it is still an AI. It can read, create, and modify real data in your Google account. Emails, calendar events, documents, contacts — all of it becomes accessible. And while that is exactly what makes the agent useful, it also means that any mistake the agent makes happens to your real data.

AI agents are not perfect. They can misinterpret instructions, especially after long conversations when older context gets compressed. They can take actions you did not ask for. And they can be influenced by external content through something called prompt injection, which we will explain below.

WARNING A real example: A user named Samar instructed her OpenClaw agent to only flag emails for deletion and never delete without permission. After a long conversation, OpenClaw compacted the context and silently dropped that instruction. The agent then deleted over 200 emails on its own.

Why Email Deserves Extra Caution

Out of all the Google services, Gmail is the one that carries the most risk. The reason is simple: your inbox is open to the outside world. Anyone can send you an email, and if your agent is reading your inbox, it will process whatever lands there. This creates an attack surface called prompt injection through email content.

Here is how it works. An attacker sends an email to your inbox that contains hidden instructions disguised as normal text. Your agent reads the email, interprets those hidden instructions as commands, and acts on them. This could mean forwarding sensitive data, changing your agent's behaviour, or executing unwanted actions on your server. More sophisticated AI models are better at resisting this, but

no model is completely immune.

RISK **What a malicious email could do:** Trick your agent into forwarding sensitive data to an external address. Make it download harmful software to your server. Instruct it to delete files, emails, or calendar events. Override your agent's safety rules and change its behaviour.

The Solution: A Dedicated Agent Account

The approach we use in this course is simple. Instead of connecting your personal Google account with all your private emails and documents, we create a completely separate Google account that belongs to the agent. This way the agent has its own clean workspace, and your personal data stays untouched no matter what happens.

Approach	Risk Level	Notes
Your personal Gmail	HIGH	Agent sees all your emails, docs, contacts. Prompt injection risk is real.
Dedicated agent account	LOW	Agent has its own clean inbox. You forward only what you want it to see.
Personal account, read-only	MEDIUM	Agent cannot send or delete, but can still be influenced by malicious emails it reads.

How to Set Up Safely

- Step 1:** Create a new Google account dedicated to your agent (e.g. ace.myagent@gmail.com).
- Step 2:** Use this new account when you go through the GOG setup in the next lesson.
- Step 3:** If you want the agent to see certain emails, set up forwarding rules from your main account to forward only specific emails to the agent account.
- Step 4:** Keep sensitive documents, financial data, and private conversations on your personal account where the agent cannot access them.
- Step 5:** If you ever need to revoke access, go to Google Cloud Console, open your project, then APIs & Services > Credentials, find the OAuth Client ID you created for OpenClaw, and delete it. This immediately invalidates the token and cuts off your agent's access. Removing the account from the test users list alone is not enough because existing tokens may continue to work for a while.

OUR APPROACH In this course we use a dedicated agent account. This way even if something goes wrong, your personal data stays completely safe. Your agent gets its own clean workspace to operate in, and you control exactly what information it has access to.

What About Calendar and Drive?

Calendar and Drive are lower risk than email because external people cannot push content into them the way they can with email. However, if you share calendars or documents with other people, those shared items could potentially contain content that influences your agent. The dedicated account approach solves this too, because the agent's calendar and Drive start completely empty and clean.

Quick Security Checklist

- ✓ Create a dedicated Google account for your agent
- ✓ Never connect your primary personal or business Gmail directly
- ✓ Add the dedicated account as a test user in Google Cloud Console
- ✓ Use email forwarding rules if you want the agent to see specific emails
- ✓ Keep financial documents, contracts, and sensitive files on your personal account
- ✓ Review what your agent is doing regularly, especially early on
- ✓ Remember: your agent's security hardening from earlier in the course still applies